

A. Kernels

1. Show that the dimension of the feature space associated to the polynomial kernel of degree d , $K: \mathbb{R}^N \times \mathbb{R}^N \rightarrow \mathbb{R}$, $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x} \cdot \mathbf{x}' + c)^d$, with $c > 0$, is

$$\binom{N+d}{d}. \quad (1)$$

It is clear that K can be written as a linear combinations of all monomials $x_1^{k_1} \cdots x_N^{k_N}$ with $\sum_{j=1}^N k_j \leq d$. The dimension of the feature space is thus the number of such monomials, $f(N, d)$, that is the number of ways of adding N non-negative integers to obtain a sum of at most d . Note that any sum of $N-1$ integers less than or equal to d can be uniquely completely to be equal to d by adding one more term. This defines in fact a bijection between sums of $N-1$ integers with value at most d and sums of N integers with value equal to d . Thus, the number of sums of N integers exactly equal to d is $f(N-1, d)$.

Now, since a sum of N terms less than or equal to d is either equal to d or less than or equal to $d-1$,

$$f(N, d) = f(N-1, d) + f(N, d-1).$$

The result then follows by induction on $N+d$, using $f(1, 0) = f(0, 1) = 1$.

Write K in terms of kernels $k_i: (\mathbf{x}, \mathbf{x}') \mapsto (\mathbf{x} \cdot \mathbf{x}')^i$, $i \in [0, d]$. What is the weight assigned to each k_i in that expression? How does it vary as a function of c .

By the binomial identity, $K(\mathbf{x}, \mathbf{x}') = (\mathbf{x} \cdot \mathbf{x}' + c)^d = \sum_{i=0}^d \binom{d}{i} c^{d-i} (\mathbf{x} \cdot \mathbf{x}')^i$. The weight assigned to each k_i , $i \leq d$, is thus $\binom{d}{i} c^{d-i}$. Increasing c decreases the weight of k_i , particularly that of k_i s with larger i .

2. For $\alpha \geq 0$, the kernel $K_\alpha: (\mathbf{x}, \mathbf{x}') \mapsto \sum_{k=1}^N \min(|x_k|^\alpha, |x'_k|^\alpha)$ over $\mathbb{R}^N \times \mathbb{R}^N$ is used in image classification. Show that K_α is PDS. To do that, you can proceed as follows.

- (a) Use the fact that $(f, g) \mapsto \int_{t=0}^{+\infty} f(t)g(t)dt$ is an inner product over the set of measurable functions over $[0, +\infty)$ to show that $(x, x') \mapsto$

$\min(x, x')$ is a PDS kernel (hint: associate an indicator function to x and another one to x').

Observe that $\min(|u|^\alpha, |u'|^\alpha) = \int_0^{+\infty} 1_{t \in [0, |u|^\alpha]} 1_{t \in [0, |u'|^\alpha]} dt$, which shows that $(u, u') \mapsto \min(|u|^\alpha, |u'|^\alpha)$ is PDS.

- (b) Use the previous question to show that K_1 is PDS and similarly K_α with other values of α .

Since $K_\alpha(x, x') = \sum_{k=1}^N \min(|x_k|^\alpha, |x'_k|^\alpha)$, K_α is PDS as a sum of N PDS kernels.

B. Boosting

1. Assume that the main weak learner assumption of AdaBoost holds. Let h_t be the base learner selected at round t . Show that the base learner h_{t+1} selected at time t must be different from h_t .

By the weak learning assumption, there exists a hypothesis $h \in H$ whose D_{t+1} -error is less than half. Examine the empirical error of h_t for the distribution D_{t+1} . Since $Z_t = 2\sqrt{\epsilon_t(1 - \epsilon_t)}$ and $\alpha_t = \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t}$,

$$\begin{aligned} \hat{R}_{D_{t+1}}(h_t) &= \sum_{i=1}^m \frac{D_t(i) e^{-\alpha_t y_i h_t(x_i)}}{Z_t} 1_{y_i h_t(x_i) < 0} \\ &= \sum_{y_i h_t(x_i) < 0}^m \frac{D_t(i) e^{\alpha_t}}{Z_t} \\ &= \frac{e^{\alpha_t}}{Z_t} \sum_{y_i h_t(x_i) < 0}^m D_t(i) \\ &= \frac{\sqrt{\frac{1 - \epsilon_t}{\epsilon_t}}}{2\sqrt{\epsilon_t(1 - \epsilon_t)}} \epsilon_t = \frac{1}{2}. \end{aligned}$$

This shows that h_t cannot be selected at round $t + 1$.

2. Let the training sample be $S = ((x_1, y_1), \dots, (x_m, y_m))$. Suppose we wish to penalize differently errors made on x_i versus x_j . To do that, we associate some non-negative importance weight w_i to each point x_i and define the objective function $F(\alpha) = \sum_{i=1}^m w_i e^{-y_i f(x_i)}$, where $f = \sum_t \alpha_t h_t$ with the notation already used in class. Show that this function is convex and differentiable and use it to derive a boosting-type algorithm (give a clear description of the algorithm similar to that of AdaBoost presented in class).

For all α , $F(\alpha) = \sum_{i=1}^m w_i e^{-y_i f(x_i)} = \sum_{i=1}^m w_i e^{-y_i \sum_{t=1}^T \alpha_t h_t(x_i)}$, with $w_i \geq 0$. F is convex as a non-negative linear combination of convex functions and is

clearly differentiable. Applying coordinate descent to this function leads to the same algorithm as AdaBoost with the only difference that

$$D_1(i) \leftarrow \frac{w_i}{\sum_{i=1}^m w_i}. \quad (2)$$

C. Perceptron

1. The margin bound on the maximum number of updates presented in class for the perceptron algorithm was given for the special case $\eta = 1$. Consider now the general perceptron update $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \eta y_t \mathbf{x}_t$, where $\eta > 0$. With the same assumptions as for the theorem presented in class, prove a bound on the maximum number of mistakes. How does η affect the bound?

The bound is unaffected as shown by the following using the same definitions and steps as in the lecture slides:

$$\begin{aligned} M\rho &\leq \frac{\mathbf{v} \cdot \sum_{t \in I} y_t \mathbf{x}_t}{\|\mathbf{v}\|} \\ &= \frac{\mathbf{v} \cdot \sum_{t \in I} (\mathbf{w}_{t+1} - \mathbf{w}_t)/\eta}{\|\mathbf{v}\|} \quad (\text{definition of updates}) \\ &= \frac{\mathbf{v} \cdot \mathbf{w}_{T+1}}{\eta \|\mathbf{v}\|} \\ &\leq \|\mathbf{w}_{T+1}\|/\eta \quad (\text{Cauchy-Schwarz ineq.}) \\ &= \|\mathbf{w}_{t_m} + \eta y_{t_m} \mathbf{x}_{t_m}\|/\eta \quad (t_m \text{ largest } t \text{ in } I) \\ &= \left[\|\mathbf{w}_{t_m}\|^2 + \eta^2 \|\mathbf{x}_{t_m}\|^2 + \underbrace{\eta y_{t_m} \mathbf{w}_{t_m} \cdot \mathbf{x}_{t_m}}_{\leq 0} \right]^{1/2} / \eta \\ &\leq \left[\|\mathbf{w}_{t_m}\|^2 + \eta^2 R^2 \right]^{1/2} / \eta \\ &\leq \left[M\eta^2 R^2 \right]^{1/2} / \eta = \sqrt{M} R. \quad (\text{applying the same to previous } ts \text{ in } I). \end{aligned}$$

2. Suppose each input vector \mathbf{x}_t , $t \in [1, T]$, coincides with the t th unit vector of \mathbb{R}^T . How many updates does it take the perceptron algorithm to converge? How does the number of updates (or mistakes) compare with the margin bound?

Clearly, it takes T updates and leads to $\mathbf{w} = \sum_{t=1}^T y_t \mathbf{x}_t$. Let $\mathbf{u} \in \mathbb{R}^T$ be a vector of norm 1 defining a separating hyperplane, thus $y_t \mathbf{u} \cdot \mathbf{x}_t = y_t u_t \geq 0$ for all $t \in [1, T]$. To obtain the maximum margin ρ , we seek a vector \mathbf{u} maximizing the minimum of $y_t u_t$ with $y_t u_t \geq 0$ for all t and $\|\mathbf{u}\| = 1$. By symmetry, all $y_t u_t$ s are equal, thus $u_t = y_t / \sqrt{T}$ for all $t \in [1, T]$ and $\rho = 1/\sqrt{T}$. Thus, Novikoff's bound gives $R^2/\rho^2 = 1/(1/T) = T$.